

# Chapter 1

## Refinement of Indexes and B+ Trees

Let us consider issues that exist in any index structure (not just B+ trees):

- Query Support: what class of queries does the index allow?
- Choice of Search Key: affects queries for which we can use an index.
- Data Entry Storage: affects performance of the index
- Variable-length key tricks: affects performance of the index.
- Cost model for index vs. heap vs. sorted file

### 1.1 Query Support

Usually indexes are used for basic select queries, including equality selection (operation is '=') and range selection (operation is one of '<', '<=', '>', '>=', 'BETWEEN'). B+ trees provide queries for both of these on a single dimension. There are other, more exotic, selections. There are queries for 2-d boxes or circles, for n-dimensional indexes a la R-trees and KD-trees. These are hard to perform well when we get to lots of dimensions ("Curse of Dimensionality" - will be reviewed at a later time). There are also near-neighbor queries ("10 restaurants closest to the Empire State Building."). There are also regular expression matches, genome string matches, and more.